

Thread-safe histograms

Jim Pivarski

Princeton University – DIANA

June 1, 2017

DAQ makes thousands of histograms.

To safely fill them in parallel threads, they're scattered into thread-local copies, gathered by histogram-addition.

But this uses so much working memory, some tasks can't be performed.

What is the cost of a thread-safe histogram?

Idea #1: make each histogram an actor with a thread-safe queue as a mailbox. “Fill” sends a message.

Idea #2: block write access at the granularity of *bins*.

using locks?

compare-and-swap?

`std::atomic?`

Idea #1: make each histogram an actor with a thread-safe queue as a mailbox. “Fill” sends a message.

No. These things work when load is variable, allowing the receiver to catch up on old messages during low load without slowing down the sender. Doing this in a steady-load environment doesn't make sense.

Idea #2: block write access at the granularity of *bins*.
using locks?

compare-and-swap?

`std::atomic?`

Idea #1: make each histogram an actor with a thread-safe queue as a mailbox. “Fill” sends a message.

No. These things work when load is variable, allowing the receiver to catch up on old messages during low load without slowing down the sender. Doing this in a steady-load environment doesn't make sense.

Idea #2: block write access at the granularity of *bins*.

using locks?

No. The memory overhead of a lock per bin would defeat the purpose (saving memory), and a lock per histogram is too coarse.

compare-and-swap?

`std::atomic`?

Idea #1: make each histogram an actor with a thread-safe queue as a mailbox. “Fill” sends a message.

No. These things work when load is variable, allowing the receiver to catch up on old messages during low load without slowing down the sender. Doing this in a steady-load environment doesn't make sense.

Idea #2: block write access at the granularity of *bins*.

using locks?

No. The memory overhead of a lock per bin would defeat the purpose (saving memory), and a lock per histogram is too coarse.

compare-and-swap?

Sounds good: per-bin contention is low and the work to be repeated in the case of a collision is minimal. No extra memory required.

`std::atomic`?

Idea #1: make each histogram an actor with a thread-safe queue as a mailbox. “Fill” sends a message.

No. These things work when load is variable, allowing the receiver to catch up on old messages during low load without slowing down the sender. Doing this in a steady-load environment doesn't make sense.

Idea #2: block write access at the granularity of *bins*.

using locks?

No. The memory overhead of a lock per bin would defeat the purpose (saving memory), and a lock per histogram is too coarse.

compare-and-swap?

Sounds good: per-bin contention is low and the work to be repeated in the case of a collision is minimal. No extra memory required.

`std::atomic`?

I originally thought this did locking, but I was wrong.

Python script sets up many threads, feeds them all the *same* block of memory (`fillme`), and starts them all at the same time.

```
double fill(long *fillme, long size, long trials, long cardinality, long *collisions) {
    struct timeval startTime, endTime;
    int shift = (int)floor(log2((double)size / cardinality)); // control collision rate

    std::mt19937 rng; // thread-local random number generator
    rng.seed(std::random_device()());
    std::uniform_int_distribution<long> distribution(0, cardinality - 1);

    gettimeofday(&startTime, 0); // start the stopwatch
    for (long i = 0; i < trials; i++) {
        long value = distribution(rng) << shift; // drop LSBs to get more collisions

        fillme[value]++; // naive increment bin
    }
    gettimeofday(&endTime, 0); // stop the stopwatch and return time
    return (1000L * 1000L * (endTime.tv_sec - startTime.tv_sec) +
            (endTime.tv_usec - startTime.tv_usec)) / 1000.0 / 1000.0;
}
```

Python script sets up many threads, feeds them all the *same* block of memory (`fillme`), and starts them all at the same time.

```
double fill(long *fillme, long size, long trials, long cardinality, long *collisions) {
    struct timeval startTime, endTime;
    int shift = (int)floor(log2((double)size / cardinality)); // control collision rate

    std::mt19937 rng; // thread-local random number generator
    rng.seed(std::random_device()());
    std::uniform_int_distribution<long> distribution(0, cardinality - 1);

    gettimeofday(&startTime, 0); // start the stopwatch
    for (long i = 0; i < trials; i++) {
        long value = distribution(rng) << shift; // drop LSBs to get more collisions

        long *ptr = &fillme[value];
        long oldval = *ptr; // try...
        long newval = oldval + 1;
        while (CAS(ptr, oldval, newval) != oldval) {
            oldval = *ptr; // ...try again
            newval = oldval + 1;
            (*collisions)++; // measure actual collision rate
        }
    }
    gettimeofday(&endTime, 0); // stop the stopwatch and return time
    return (1000L * 1000L * (endTime.tv_sec - startTime.tv_sec) +
            (endTime.tv_usec - startTime.tv_usec)) / 1000.0 / 1000.0;
}
```

Python script sets up many threads, feeds them all the *same* block of memory (fillme), and starts them all at the same time.

```
double fill(long *fillme, long size, long trials, long cardinality, long *collisions) {
    struct timeval startTime, endTime;
    int shift = (int)floor(log2((double)size / cardinality));    // control collision rate

    // reinterpret the block of memory as an array of atomics
    std::atomic<long>* fillme2 = reinterpret_cast<std::atomic<long>*>(fillme);


    std::mt19937 rng;                                           // thread-local random number generator
    rng.seed(std::random_device()());
    std::uniform_int_distribution<long> distribution(0, cardinality - 1);

    gettimeofday(&startTime, 0);                                // start the stopwatch
    for (long i = 0; i < trials; i++) {
        long value = distribution(rng) << shift;                // drop LSBs to get more collisions

        // fancy fill method
        fillme2[value].fetch_add(1, std::memory_order_relaxed);

    }
    gettimeofday(&endTime, 0);                                    // stop the stopwatch and return time
    return (1000L * 1000L * (endTime.tv_sec - startTime.tv_sec) +
            (endTime.tv_usec - startTime.tv_usec)) / 1000.0 / 1000.0;
}
```

- ▶ Knight's Landing (for 128 threads), memory block in normal DRAM.
- ▶ Number of threads and naive vs. safe performed in a random order.
- ▶ Before execution, threads pinned to a random subset of CPUs.
- ▶ Collision rate controlled by bit shift, then measured in situ.



Compare-and-swap results

